

Do AI models of DNA actually understand biology?

MINA: linear evaluation reveals coding-sequence family signal in frozen DNA encoders up to protein composition

Austin Senna Wijaya* Hayden Leung* Department of Computer Science, Columbia University **both presenting*

MLCB 2026 spotlight

QR
repo /
paper
(pending)

The question

DNA “foundation models” are increasingly used as off-the-shelf feature extractors, and often treated as interchangeable. **MINA asks a simple question:** what do their frozen embeddings *already* know, with no extra training — real biology, or surface patterns?

How we tested it

3,244 human genes in 5 families

● TF ● GPCR ● Kinase
● Ion channel ● Immune receptor

4 frozen DNA encoders

DNABERT-2 · NT-v2 ·
GENA-LM · HyenaDNA
weights never trained →
a fixed list of numbers

Simple linear probe
the only part that learns

Two read-outs

Q1 · which protein family?
Q2 · predict the gene’s text description

Benchmarked against **counting amino acids** (composition) and
a **real protein model** (ESM-2)

Words to know

Encoder — an AI model that turns a DNA sequence into a list of numbers.


Frozen — the model is used as-is; we never retrain it.

Linear probe — a tiny read-out we do train, to test what the numbers already contain.

CDS — a gene’s protein-coding sequence. **TSS window** — the region around a gene’s start, mostly regulatory (non-coding).


Homology-aware split — keep near-identical genes on the same side of train/test, so the model can’t just recognise a look-alike.

Coding genes carry a family signal



From a gene’s **coding sequence**, the families separate on their own (GPCRs, orange, split off cleanly). Feed the same genes’ **regulatory window** and the structure dissolves into one cloud.


... but only up to composition



The best DNA encoder (NT-v2, **0.727**) beats simple 4-mer composition (0.633) — but only *ties* translated amino-acid composition (0.735), and sits far below a dedicated protein model, ESM-2 (**0.960**). And encoders are *not* interchangeable: GENA-LM (0.415) falls below the baseline.


Paired test: vs. 4-mer +0.094 [+0.030, +0.166]; vs. amino-acid composition −0.008 [−0.064, +0.056], a statistical tie.

Regulatory context collapses



Swap coding sequence for the gene’s 196,608 bp regulatory (TSS) window and family signal falls to near chance (NT-v2 **0.313** vs. 0.224). Predicting a gene’s natural-language description is weak everywhere (best $R^2 = 0.077$ on coding sequence; near zero on the TSS window). The same collapse holds under a stricter *disjoint*-TSS split (NT-v2 0.313 → 0.259).

Honest evaluation matters



A naive **random** split inflates scores by letting near-identical genes leak between train and test. ESM-2 barely moves (0.96); the DNA encoders drop most (NT-v2 0.83 → 0.73), because part of their apparent skill was leaked look-alikes.

So what? Don’t treat DNA encoders as generic gene-function models. Trust a benchmark only if it has **(1)** a translated-composition baseline and **(2)** a homology-aware split — otherwise memorised similarity masquerades as understanding. **MINA is a cheap diagnostic that tells the two apart.**